深度學習智慧應用
Lab 4: FCN、U-Net and TrackNet

Yu Ching Lin
Network Optimization Lab (NOL)
Department of Computer Science
National Chiao Tung University

Outline

- Heatmap
- Semantic Segmentation
- FCN
- U-Net
- TrackNet

HeatMap

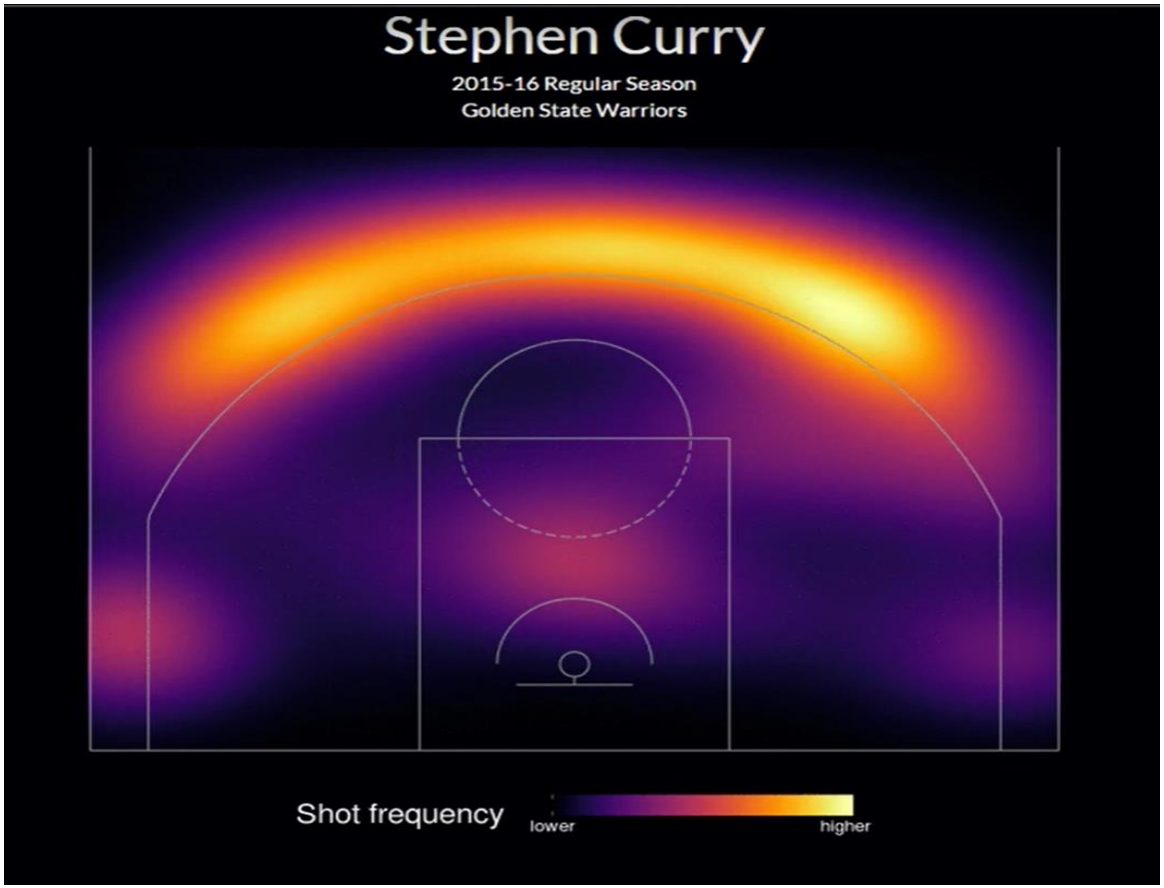
Stephen Curry
2015-16 Regular Season
Golden State Warriors
Shot frequency
lower
higher

Semantic Segmentation

Pixel-wise classification

FCN

Convolutionalization： 全連接層->卷積層

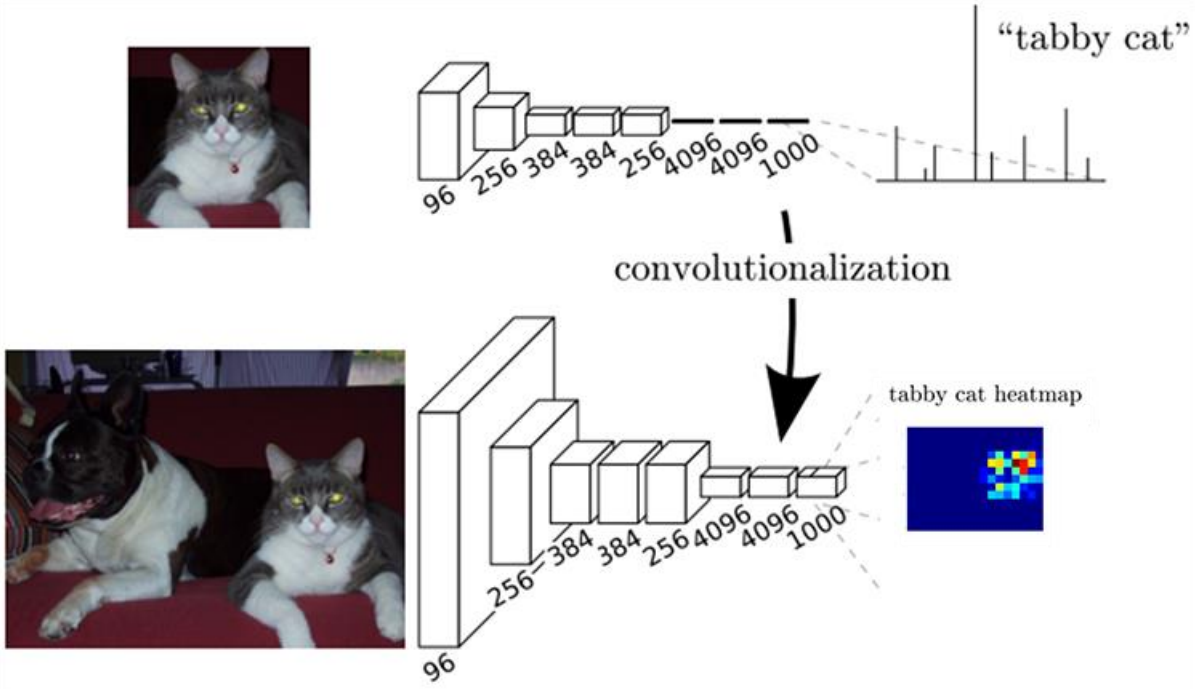

FCN

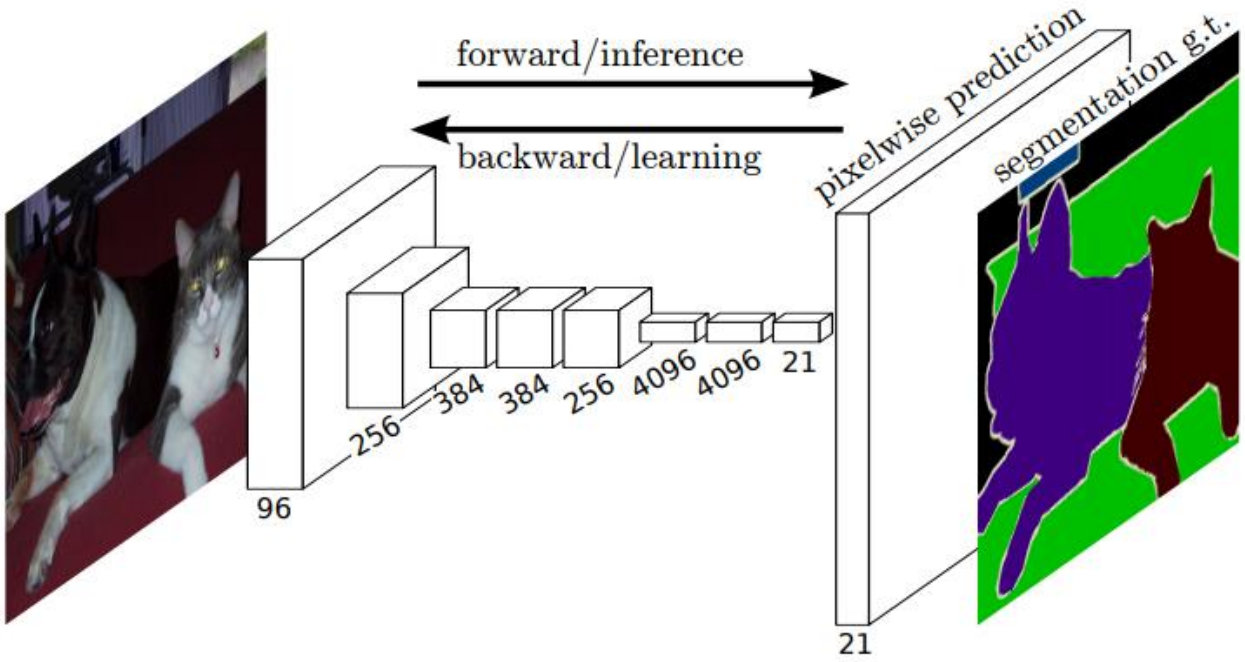
forward/inference
backward/learning
pixelwise prediction
segmentation g.t.
96
256
384
384
256
4096
4096
21
21

FCN

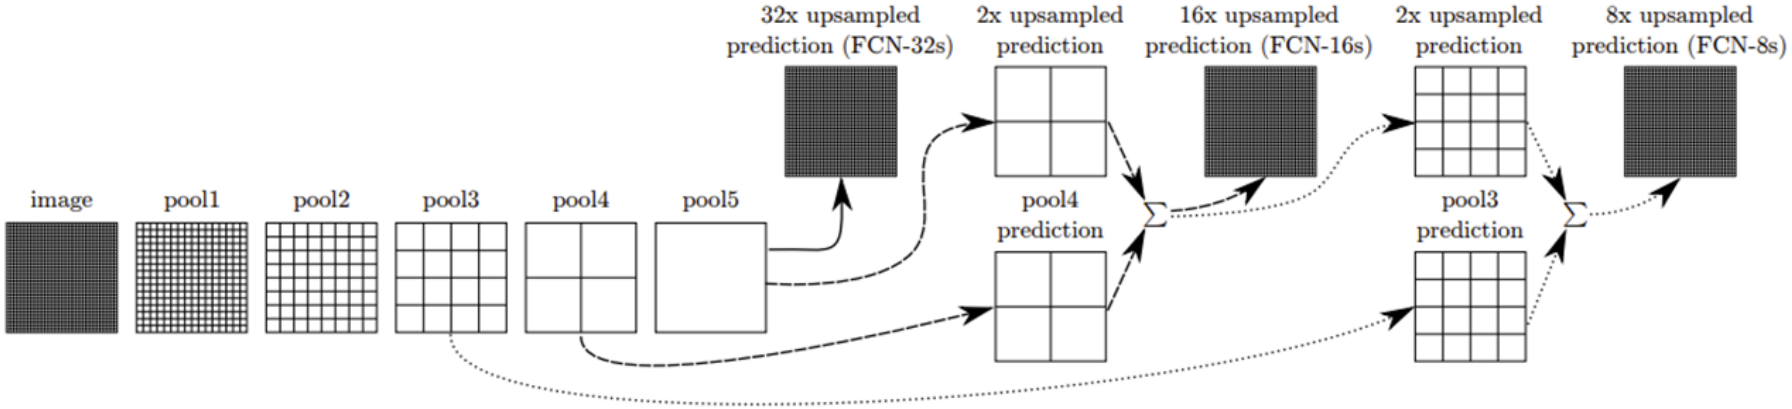
32x upsampled prediction (FCN-32s)
2x upsampled prediction
16x upsampled prediction (FCN-16s)
2x upsampled prediction
8x upsampled prediction (FCN-8s)
image
pool1
pool2
pool3
pool4
pool5
pool4 prediction
pool3 prediction

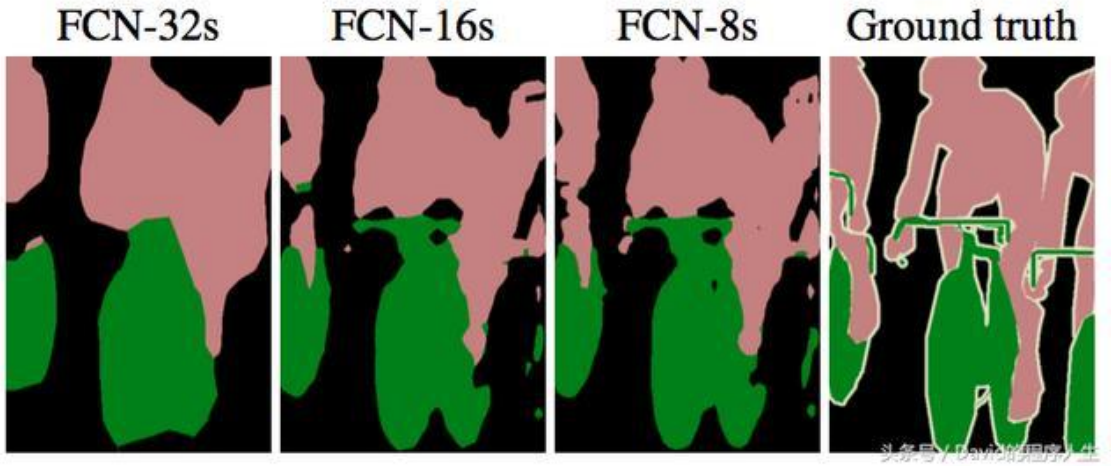
FCN-32s
FCN-16s
FCN-8s
Ground truth

U-Net

Modify and extend the architecture of fully convolutional network such that it works with very few training images and yields more precise segmentations

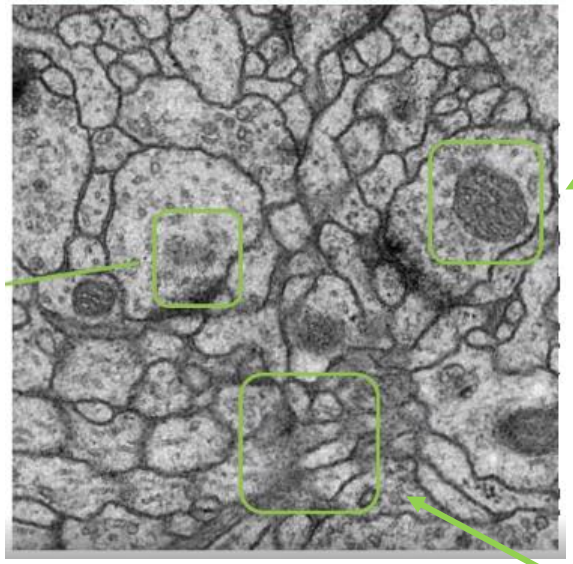

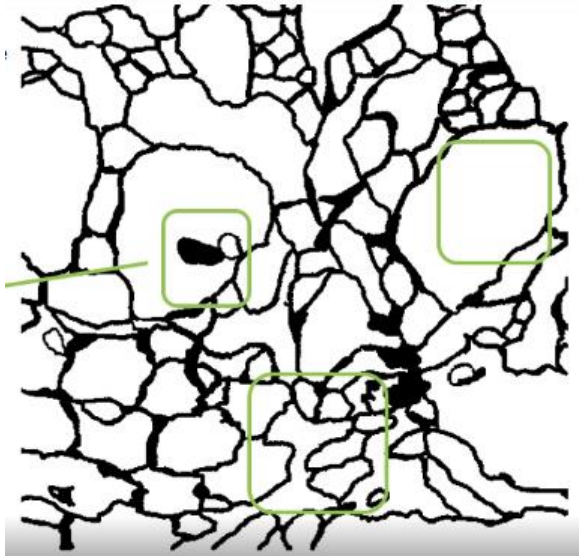

U-Net Architecture

Concatenation with the correspondingly cropped feature map from the contracting path

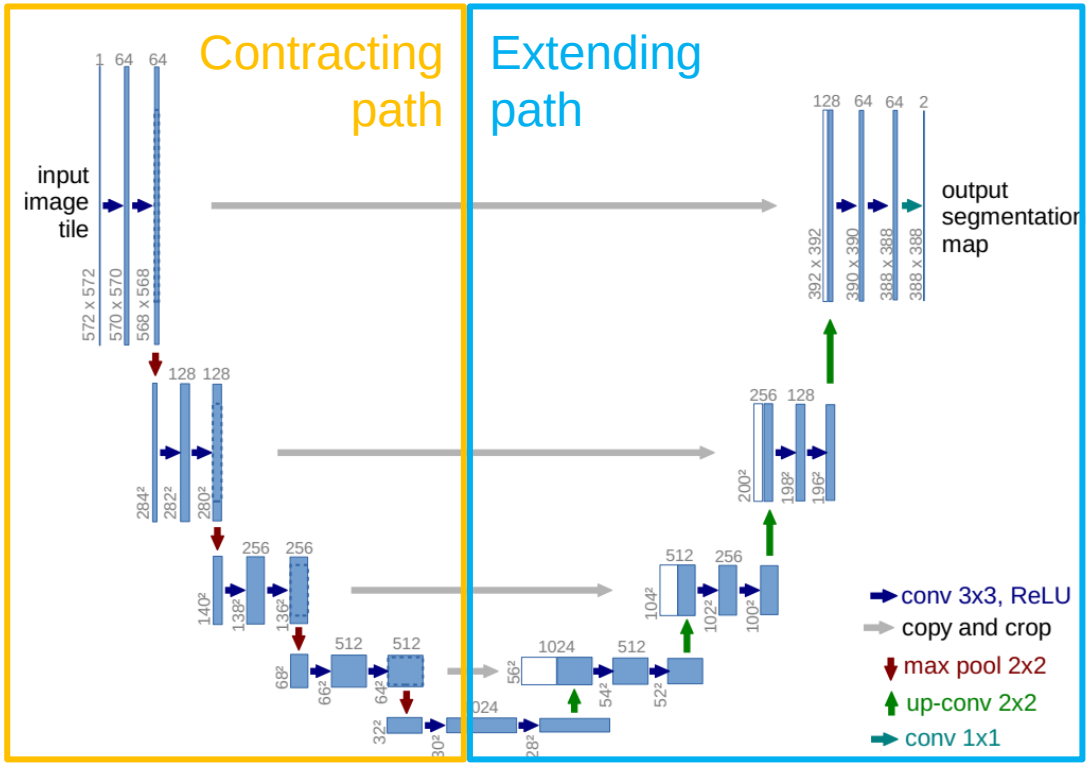

TrackNet

https://www.pyimagesearch.com/2018/10/29/multi-object-tracking-with-dlib/

Multiple Object Tracking

Single Object Tracking

https://sandipanweb.wordpress.com/category/computer-vision/

https://nanonets.com/blog/object-tracking-deepsort/

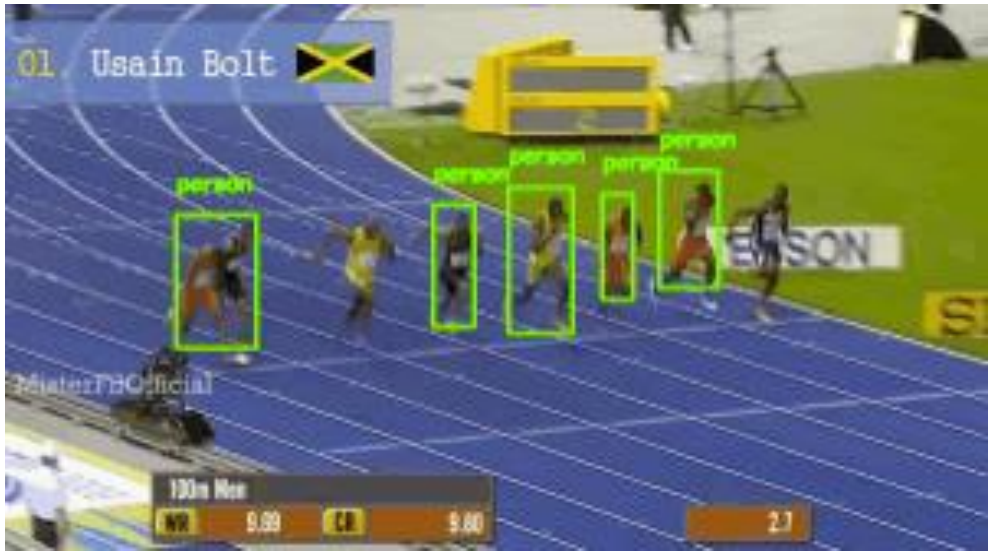

TrackNet

https://people.cs.nctu.edu.tw/~yi/TechReports/TrackNet.v1_Final_NOL.pdf

標記網球位置的時，網球呈現軌跡拖延的案例

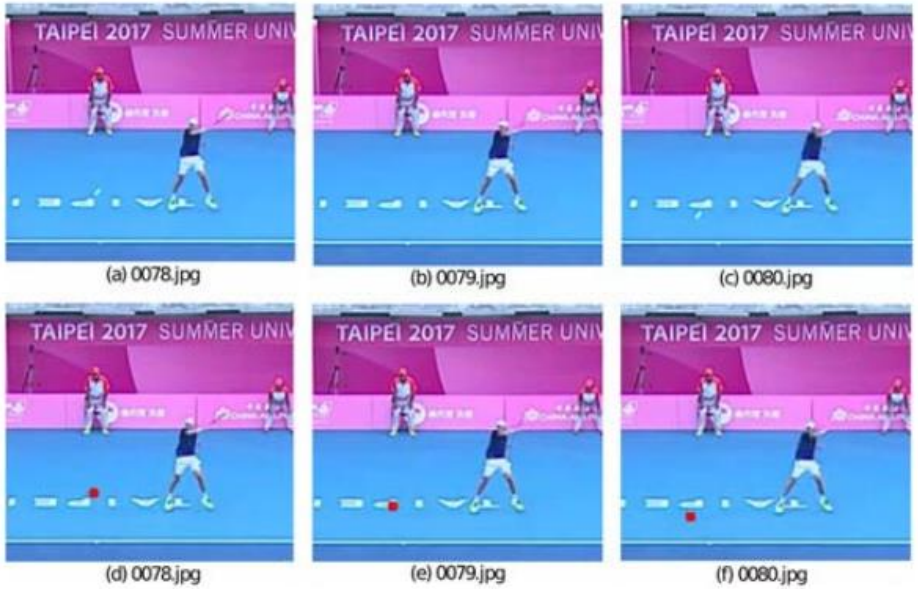

網球影像無法被識別

TrackNet

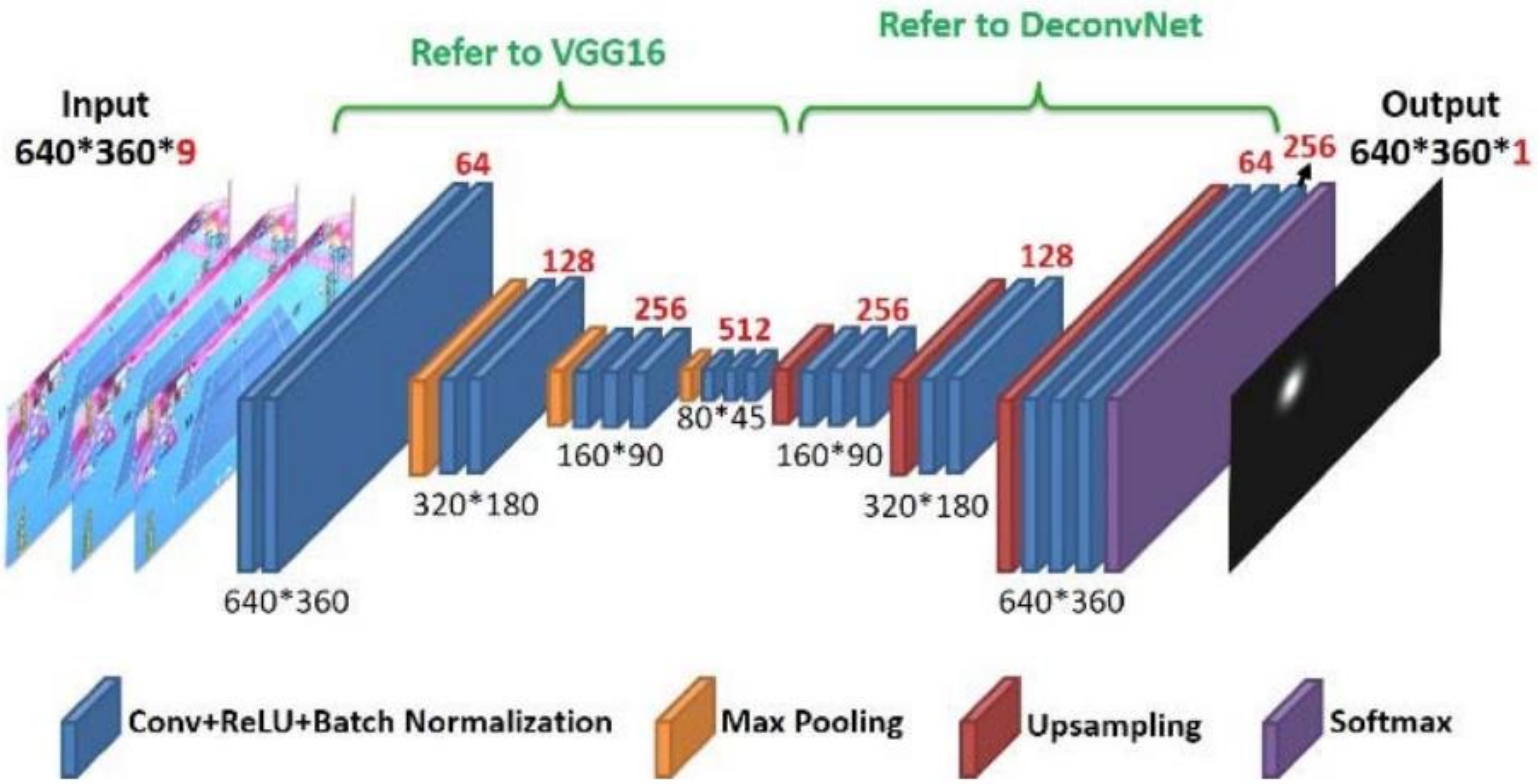
Refer to VGG16
Refer to DeconvNet
Input
640*360*9
Output
640*360*1
64
128
256
512
256
128
64
256
640*360
320*180
160*90
80*45
160*90
320*180
640*360
Conv+ReLU+Batch Normalization
Max Pooling
Upsampling
Softmax

TrackNet

- Performance

Lab 4

Outline

- Data Label
- Data Generate
- Model Architecture
- Evaluation standard
- Ways to improve
- Score

Data Label

Label Tool

Just follow the step, you will get label result as csv file
HackMD : https://hackmd.io/@JVTTHln7SmqjN39abZ9DLQ/B1lpZCBPj

Rules:

1. Please mark the ball head of shuttlecock (follow forward direction)

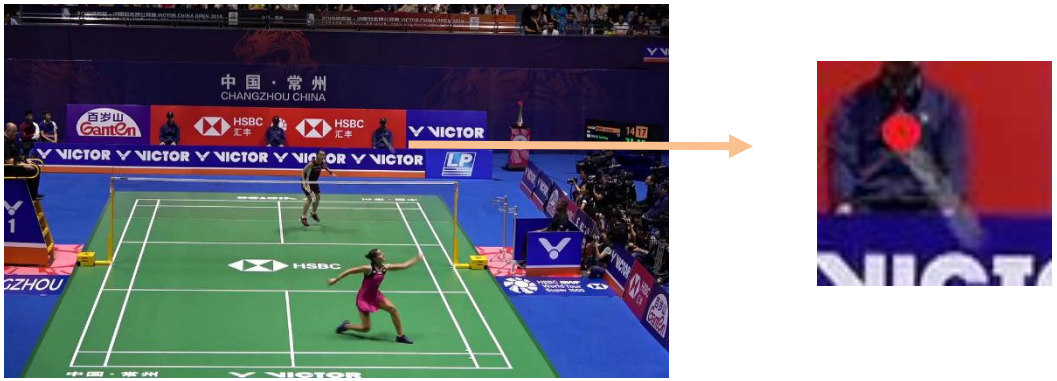

2. If you can't see the ball, skip that.

Input : mp4, output : csv

Data Label

學號/姓名	Match	video_name_1	video_name_2
[redacted]	1	1_01_00.mp4	1_01_01.mp4
[redacted]	1	1_01_03.mp4	1_01_04.mp4
[redacted]	1	1_02_05.mp4	1_02_06.mp4
[redacted]	1	1_02_08.mp4	1_02_09.mp4
[redacted]	1	1_02_10.mp4	1_03_11.mp4

- Label list :
https://docs.google.com/spreadsheets/d/1V2BTacTkfsGQEgjU_1Yfs4c4a_7GOKWzLELAcjEaJJs/edit#gid=1385358667
- Data :
https://drive.google.com/drive/folders/1wrODR97BXw9Rd3uPY-bQVwl0s5K0I1Pr?usp=share_link
- Everyone need to label 2 videos

- Upload your csv file (after labeling) in the same match folder

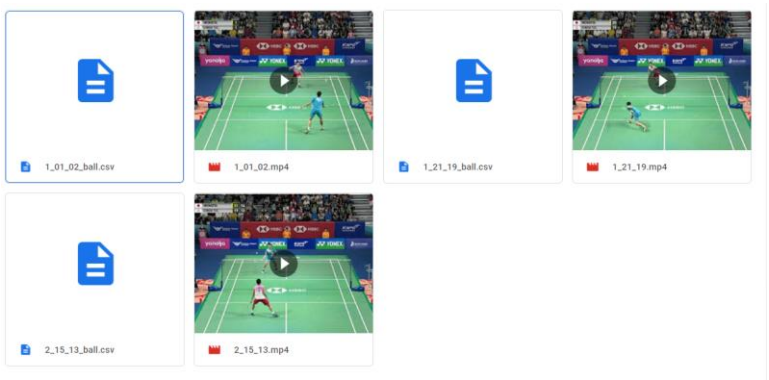

Data Generate

• 資料結構

- 將Label完的csv檔放到對應的比賽資料夾下方

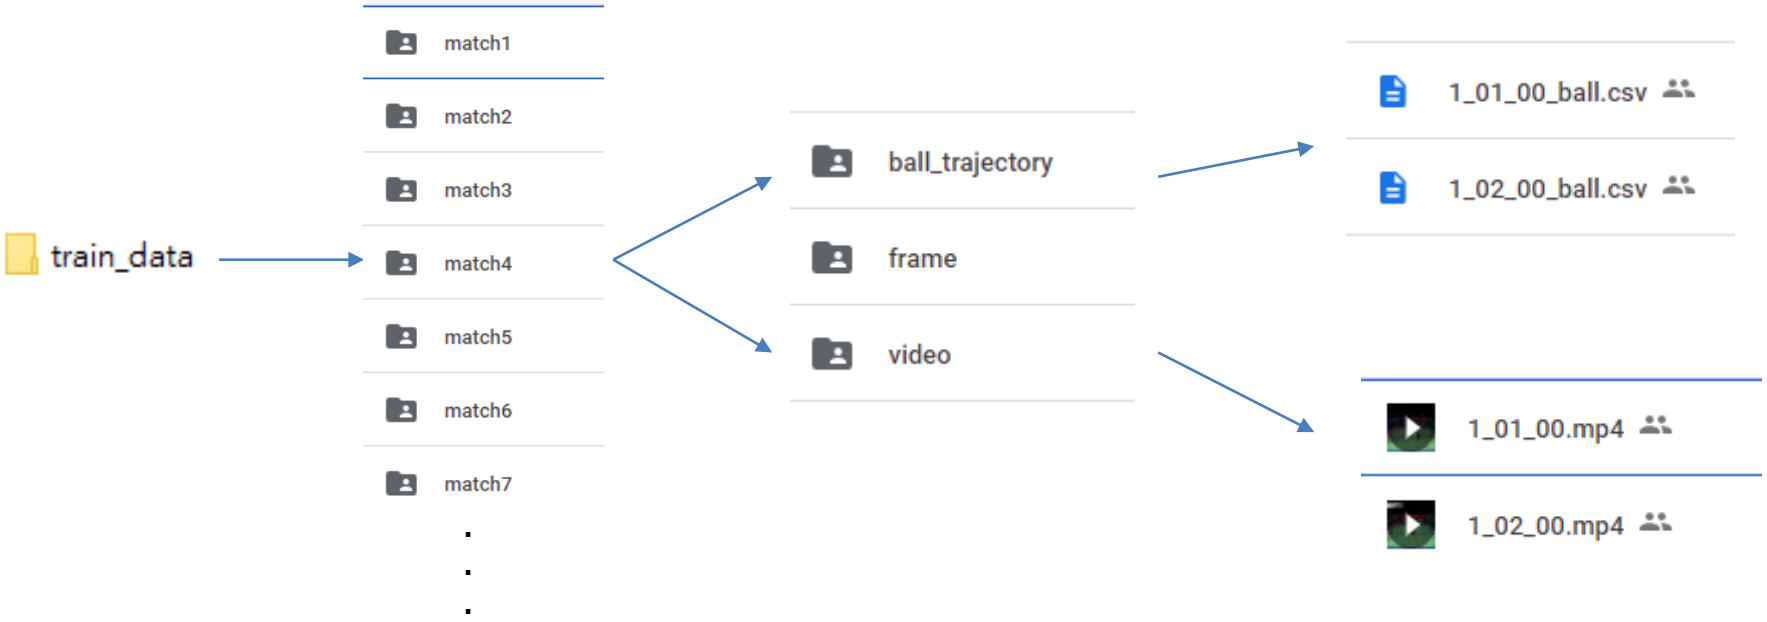

• run_frame_gen.ipynb

- run_frame_gen.ipynb

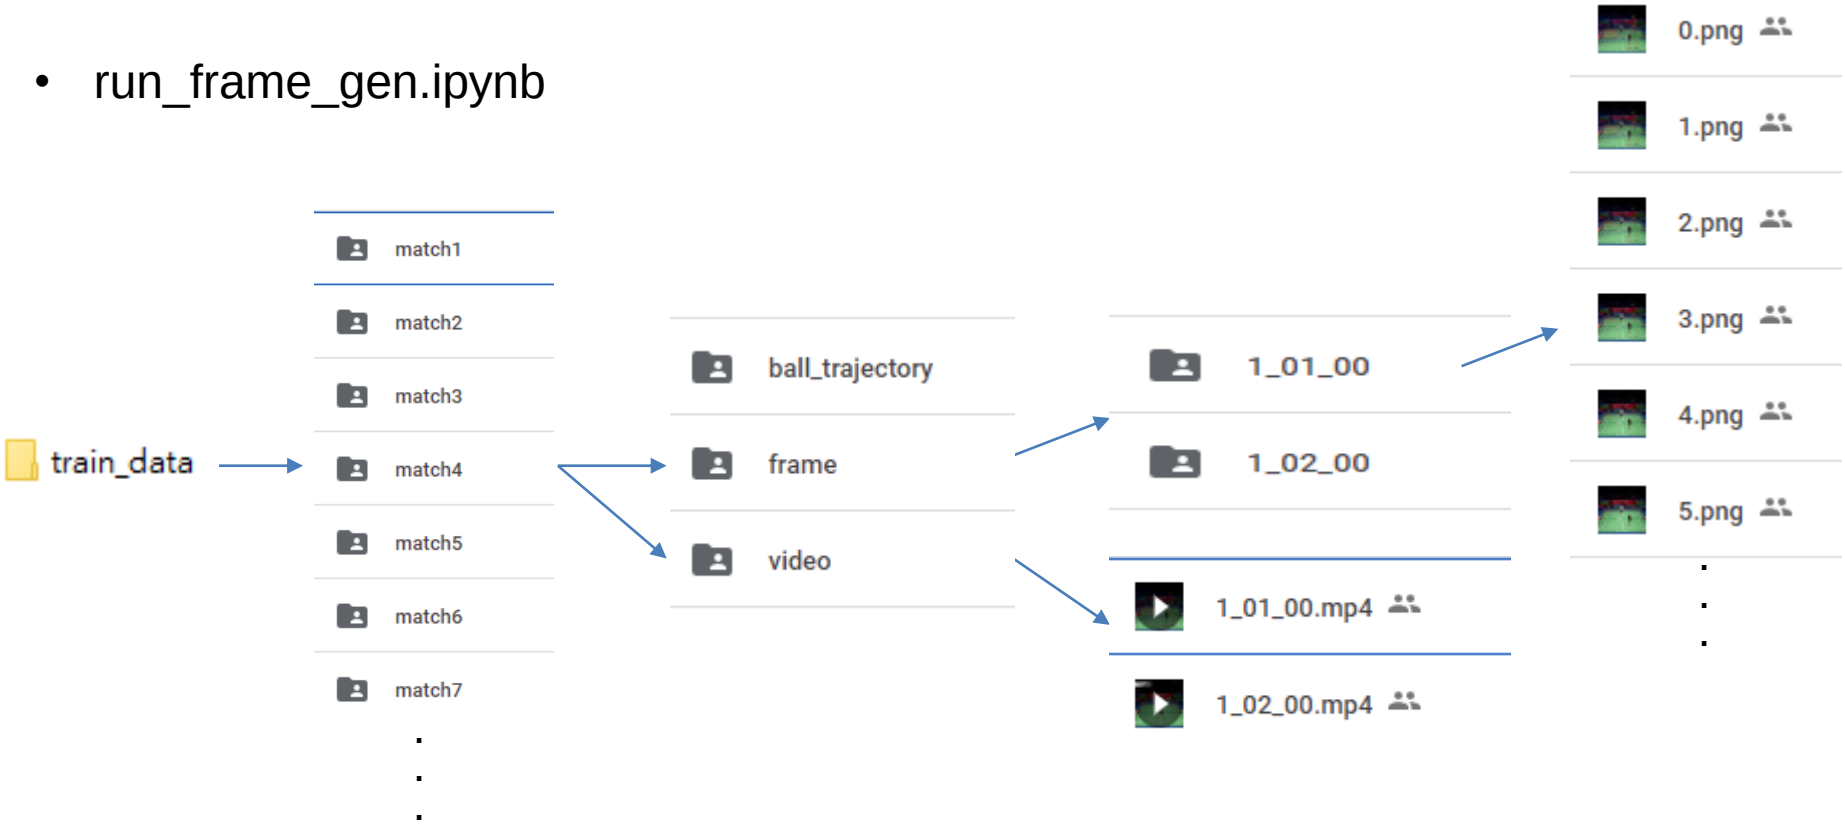

- 將train_data底下的所有比賽影片切成一個一個frame存起來

• run_gen_data.ipynb

- run_gen_data.ipynb

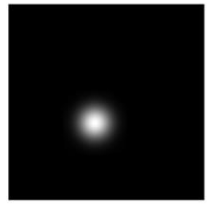

$\hat{\mathbf{Y}} \in (0,1)^{H \times W}$

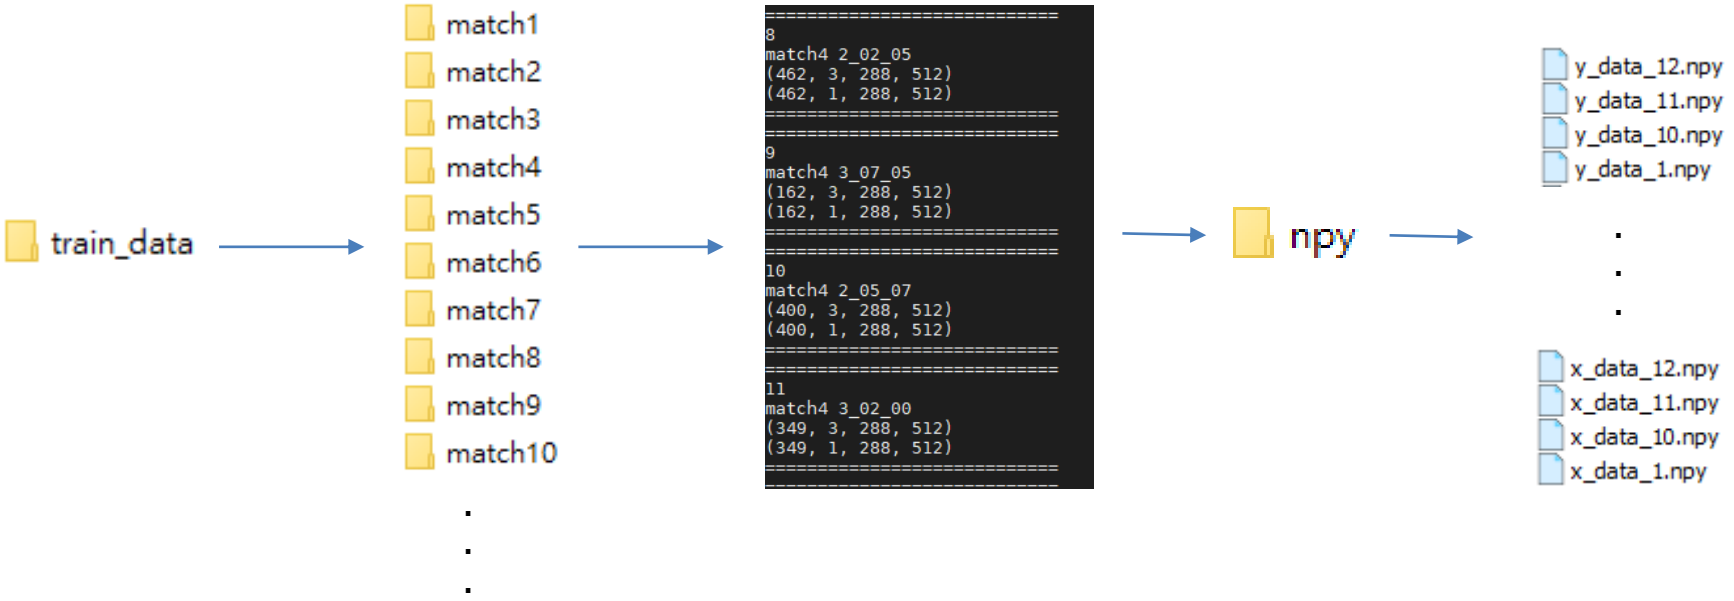

Customizing batch size

- 由於Colab的GPU記憶體空間不足，預設的batch size皆為100，若想調整batch size，更改的地方如下。

```
1   import numpy as np
2   import os
3   from glob import glob
4   import piexif
5   from keras.preprocessing.image import ImageDataGenerator
6   from tensorflow.keras.utils import array_to_img, img_to_array, load_img
7   import pandas as pd
8   from sklearn.model_selection import train_test_split
9   from keras.models import *
10  from keras.layers import *
11  import keras.backend as K
12  from keras import optimizers
13  import tensorflow as tf
14  import random
15  import shutil
16  BATCH_SIZE=100
17  HEIGHT=288
18  WIDTH=512
19  mag = 1
20  sigma = 2.5
```

gen_data_rally.py

Model architecture

TrackNet.py

Modify model by yourself ! (If you need)

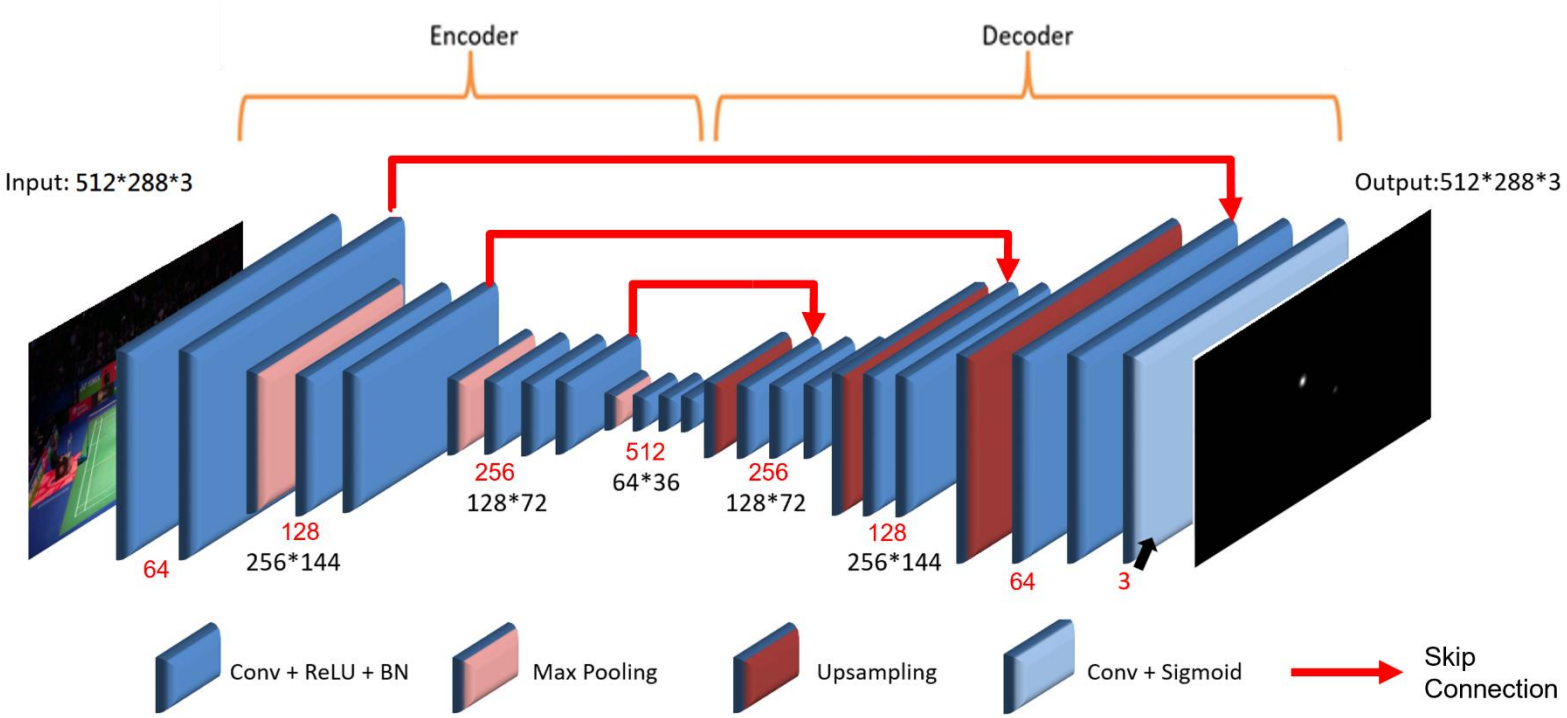

Base model

Loss function

Binary cross entropy

y = ground truth
$\hat{y}$ = predict result

$$BCE = -\frac{1}{N}\sum_{i=0}^{N} y_i \cdot log(\hat{y}_i) + (1 - y_i) \cdot log(1 - \hat{y}_i)$$

run_train_TrackNetv2.ipynb

- Colab選擇GPU環境

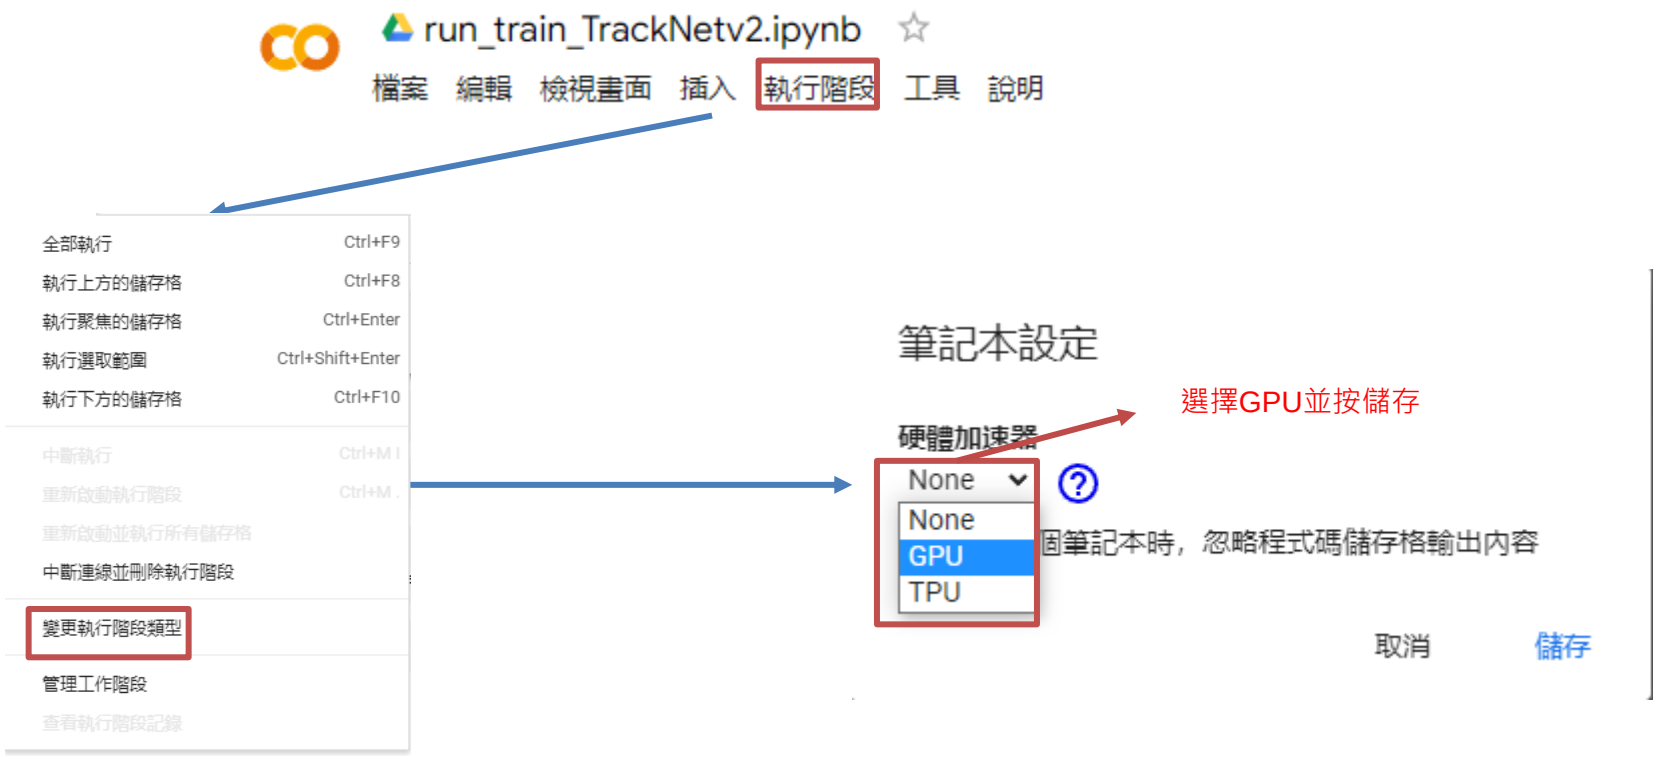

run_train_TrackNetv2.ipynb

- python3 train_TrackNet.py –load_weights = **<previousWeightPath>** --save_weights = **<weightPath>** --dataDir = **<npyDataDirectory>** --epochs = **<trainingEpochs>** --tol = **<toleranceValue>**

  - **<previousWeightPath>** is TrackNet weight you had trained before.
  - **<weightPath>** is TrackNet weight after this training
  - **<npyDataDirectory>** is the directory of the .npy training data
  - **<toleranceValue>** means tolerance value of true positive (we use 5)

EX :

- 從頭開始train

python3 train_TrackNet3.py --save_weights=mymodel --dataDir=npy --epochs=30 --tol=5

- 加train

python3 train_TrackNet3.py --load_weights=mymodel/mymodel_3 --save_weights=mymodel --dataDir=npy --epochs=30 --tol=5

Evaluation standard

Metric

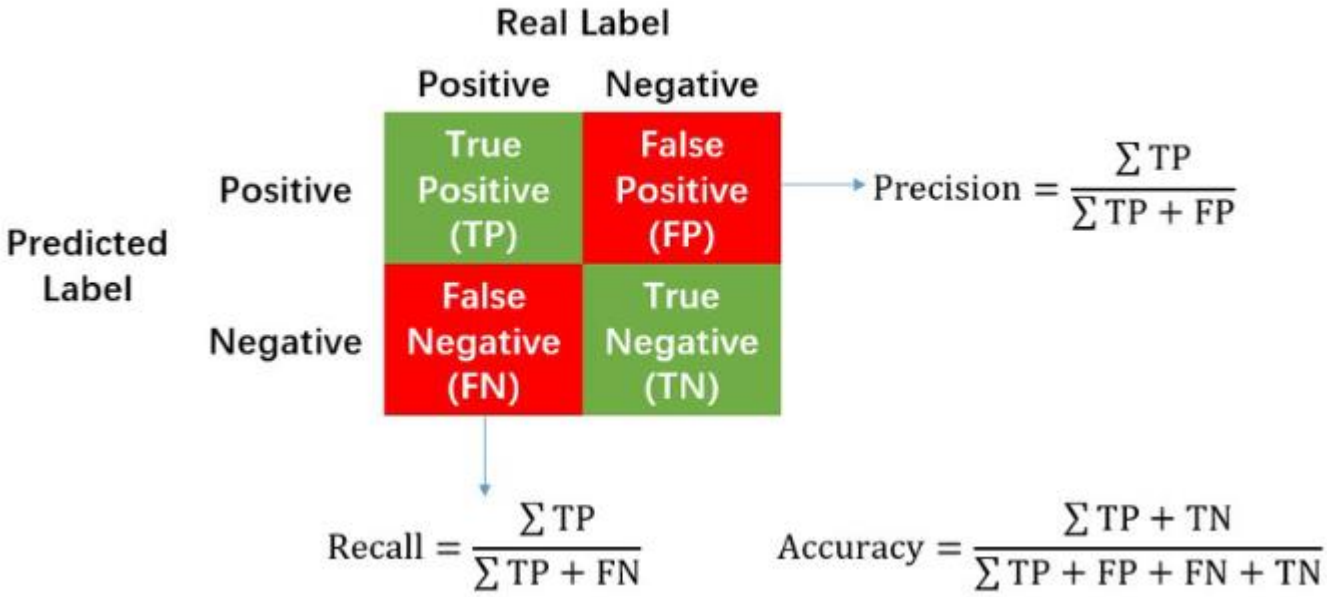

- FP1 :

Both prediction and ground truth are ball existing, but the distance is out of tolerance value.

- FP2 :

The prediction is ball existing, but the ground truth is no ball

run_predict.ipynb

- Colab也一樣要先選GPU環境
- python3 predict.py --video_name = **<videoPath>** --load_weights = **<weightPath>** --label = **<csvFile>** --tol = **<toleranceValue>**
- **<videoPath>** is video you want to predict.
- **<weightPath>** is weight of model
- **<csvFile>** is the ground truth csv of video
- **<toleranceValue>** means tolerance value of true positive (we use 5)

EX:

python predict.py --video_name=1_02_03.mp4 --load_weights=model_27 --label=1_02_03_ball.csv --tol=5

run_predict.ipynb

```
Beginning predicting......
==========================================
Number of true positive: 411
Number of true negative: 30
Number of false positive FP1: 188
Number of false positive FP2: 54
Number of false negative: 33
accuracy: 0.6159217877094972
precision: 0.6294027565084227
recall: 0.9256756756756757
(ACC. + Pre. + Rec.)/3: 0.7236667399645319
```

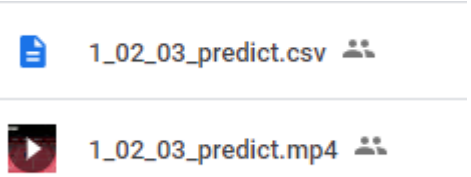

Ways to improve

Model Architecture

- Multiple Input (or output)
- Channel (Ex: 32 -> 64, 64 -> 128 …)
- Skip connection (while upsampling, like Unet)
- Drop out layer (If overfitting)
- Different activation function

It may cause GPU out of memory …

Multiple input (or output)

The devil is in the details

gen_data_rally.py

```
while ptr < num:
    x_data_tmp = []
    y_data_tmp = []
    while (i < ptr+BATCH_SIZE) and (i < num):
        if no[i]+2 != no[i+2]:
            i += 1
            continue
        unit = []
        for j in range(1):
            target=str(no[i+j])+'.png'
            png_path = os.path.join(r, target)
            a = load_img(png_path)
            a = np.moveaxis(img_to_array(a.resize(size=(WIDTH, HEIGHT))), -1, 0)
            unit.append(a[0])
            unit.append(a[1])
            unit.append(a[2])
            del a
        x_data_tmp.append(unit)
        del unit

        unit = []
        for j in range(1):
            if v[i+j] == 0:
                unit.append(genHeatMap(WIDTH, HEIGHT, -1, -1, sigma, mag))
            else:
                unit.append(genHeatMap(WIDTH, HEIGHT, int(x[i+j]/ratio), int(y[i+j]/ratio), sigma, mag))
        y_data_tmp.append(unit)
        del unit
```

1-in-1-out

```
while ptr < num-2:
    x_data_tmp = []
    y_data_tmp = []
    while (i < ptr+BATCH_SIZE) and (i < num-2):
        if no[i]+2 != no[i+2]:
            i += 1
            continue
        unit = []
        for j in range(3):
            target=str(no[i+j])+'.png'
            png_path = os.path.join(r, target)
            a = load_img(png_path)
            a = np.moveaxis(img_to_array(a.resize(size=(WIDTH, HEIGHT))), -1, 0)
            unit.append(a[0])
            unit.append(a[1])
            unit.append(a[2])
            del a
        x_data_tmp.append(unit)
        del unit

        unit = []
        for j in range(3):
            if v[i+j] == 0:
                unit.append(genHeatMap(WIDTH, HEIGHT, -1, -1, sigma, mag))
            else:
                unit.append(genHeatMap(WIDTH, HEIGHT, int(x[i+j]/ratio), int(y[i+j]/ratio), sigma, mag))
        y_data_tmp.append(unit)
        del unit
```

3-in-3-out

Multiple input (or output)

The devil is in the details

train_TrackNet3.py

```python
def outcome(y_pred, y_true, tol):
    n = y_pred.shape[0]
    i = 0
    TP = TN = FP1 = FP2 = FN = 0
    while i < n:
        for j in range(1):
            if np.amax(y_pred[i][j]) == 0 and np.amax(y_true[i][j]) == 0:
                TN += 1
            elif np.amax(y_pred[i][j]) > 0 and np.amax(y_true[i][j]) == 0:
                FP2 += 1
            elif np.amax(y_pred[i][j]) == 0 and np.amax(y_true[i][j]) > 0:
                FN += 1
            elif np.amax(y_pred[i][j]) > 0 and np.amax(y_true[i][j]) > 0:
                h_pred = y_pred[i][j] * 255
                h_true = y_true[i][j] * 255
                h_pred = h_pred.astype('uint8')
                h_true = h_true.astype('uint8')
```

1-in-1-out

```python
def outcome(y_pred, y_true, tol):
    n = y_pred.shape[0]
    i = 0
    TP = TN = FP1 = FP2 = FN = 0
    while i < n:
        for j in range(3):
            if np.amax(y_pred[i][j]) == 0 and np.amax(y_true[i][j]) == 0:
                TN += 1
            elif np.amax(y_pred[i][j]) > 0 and np.amax(y_true[i][j]) == 0:
                FP2 += 1
            elif np.amax(y_pred[i][j]) == 0 and np.amax(y_true[i][j]) > 0:
                FN += 1
            elif np.amax(y_pred[i][j]) > 0 and np.amax(y_true[i][j]) > 0:
                h_pred = y_pred[i][j] * 255
                h_true = y_true[i][j] * 255
                h_pred = h_pred.astype('uint8')
                h_true = h_true.astype('uint8')
```

3-in-3-out

Other Loss function

- Mean square error
- α-balanced cross-entropy loss
- Focal loss
  https://github.com/umbertogriffo/focal-loss-keras/blob/master/src/loss_function/losses.py
- Weighted Hausdorff distance
  https://github.com/N0vel/weighted-hausdorff-distance-tensorflow-keras-loss/blob/master/weighted_hausdorff_loss.py

Training tips

- Learning rate
- Batch size
- Data splitting
- Visualize loss

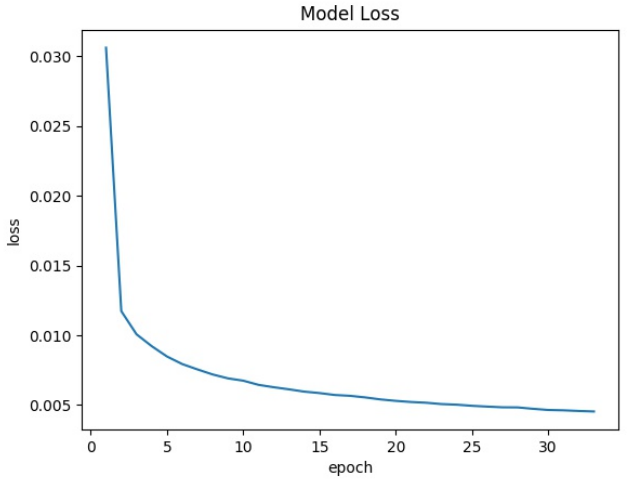

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!

Score

How we score

- Label (30%)
- Model performance (45%)

  45% for Performance benchmark: (accuracy + precision + recall) / 3

- Demo & QA(25%)

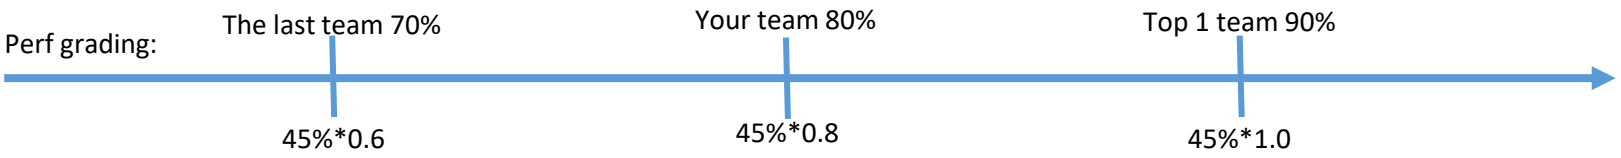

Test Data

Match not in train data

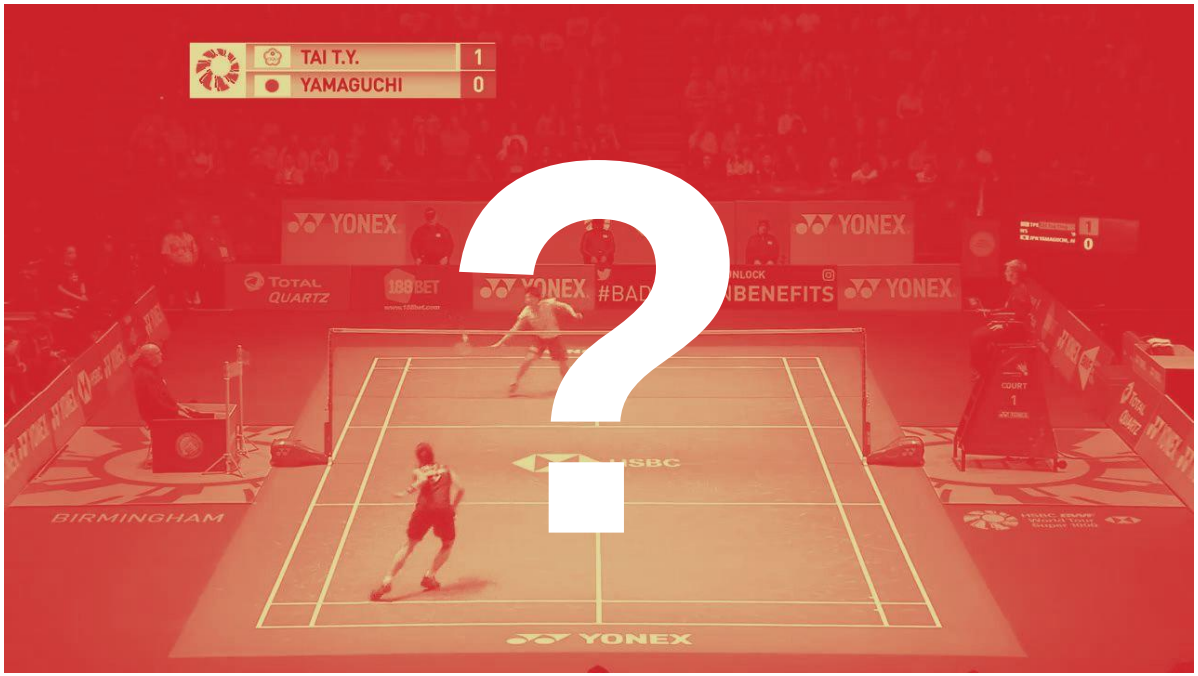

Test Data

Maybe …

Deadline

- Create data
  - Before 12/04 23:59 (half less per day, e.g. 100 → 50 → 25 …)
- Demo
  - On 12/15 after TA's lecture (about 19:30)
- Each group needs at least 1 member to demonstrate
- Lab Submission
  - Before 12/15 23:59
- If you have any problems about submission & demo please contact TAs before deadline. ([深度學習智慧應用] 主旨開頭請加這個)
  - bobhsiao0306.cs10@nycu.edu.tw
  - jason880102.cs10@nycu.edu.tw

繳交檔案格式

- lab4_group{GROUP_NUMBER}.zip
- EX：lab4_group20.zip

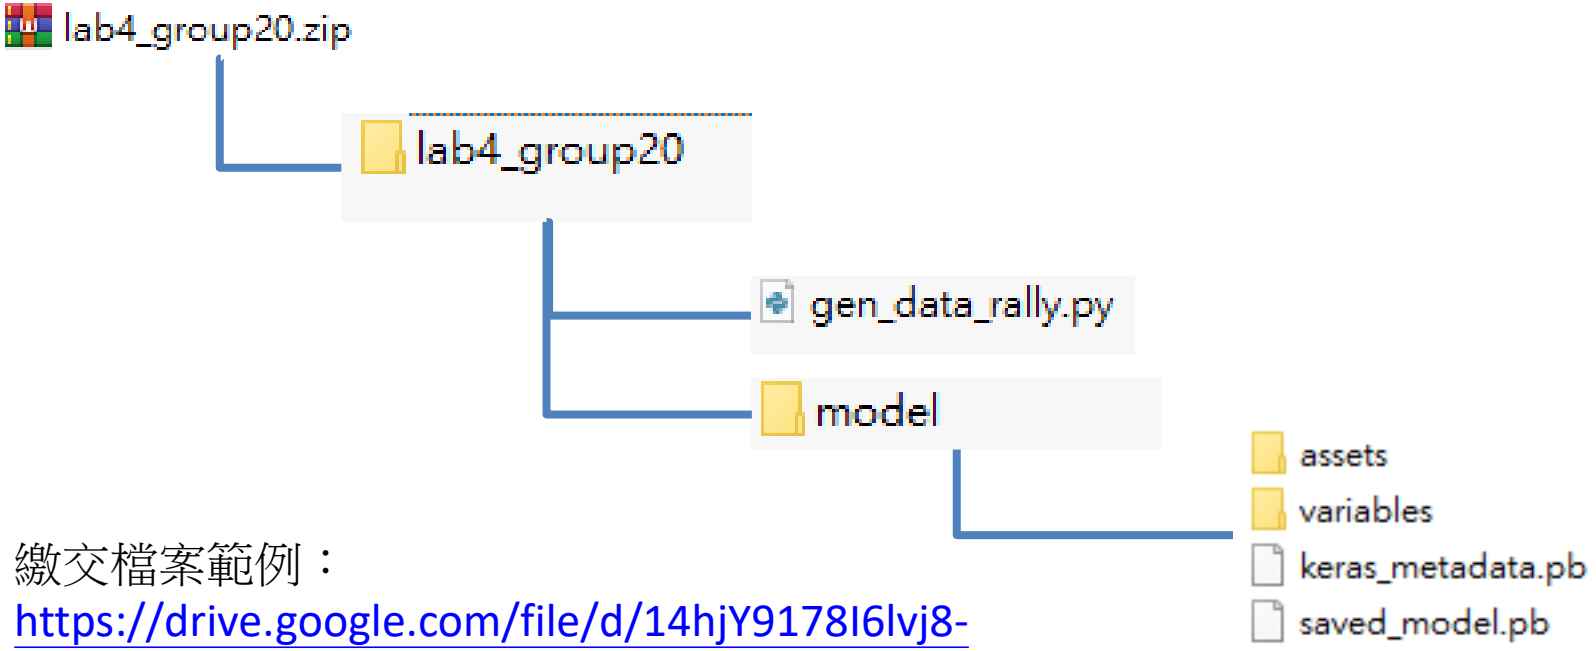

繳交檔案範例：
https://drive.google.com/file/d/14hjY9178I6lvj8-9D4Hc7FUKO65jDvwr/view?usp=sharing

Further Reading

TrackNet : A Deep Learning Network for Tracking High-speed and Tiny Objects in Sport Applications
https://people.cs.nctu.edu.tw/~yi/TechReports/TrackNet.v1_Final_NOL.pdf

High Performance Visual Tracking with Siamese Region Proposal Network (CVPR 2018)
http://www.zhengzhu.net/upload/P6938bc861e8d4583bf47d47d64ed9598.pdf

Object as points (CVPR 2019)
https://arxiv.org/pdf/1904.07850.pdf

Locating Objects Without Bounding Boxes (CVPR 2019)
https://arxiv.org/pdf/1806.07564.pdf